


Problem 4:

a.)



Linear regression will fit the training data well since the data points seem to have an almost linear trend. It will also do decent on the test data set; however since the data point are slightly spread out towards the lower and higher x's there might be a slightly larger error.

b.) Completed this portion with the following code:

```
# part g: modify to create matrix for polynomial model
X = np.append(np.ones((n, d)), X, axis=1)
```

c.) Completed this portion by using the following code:

```
y = np.dot(X, self.coef_)
```

d.)

(i) Completed this portion by using the following code:

```
# part d: compute J(theta)
n, d = X.shape
cost = 0
temp = np.dot(X, self.coef_) - y
temp *= temp
cost = temp.sum()
```

(ii) Completed this portion with the following code:

```
# hint: you cannot use self.predict(...) to make the predictions
y_pred = np.dot(self.coef_, X.T)
self.coef_ = self.coef_ - (2*eta*(np.dot(X.T, y_pred-y)))
err_list[t] = np.sum(np.power(y - y_pred, 2)) / float(n)
```

(iii)

Learning Rate	Number of iterations	Value of Objective Function
0.0001	10000	4.0863970367957645
0.001	7021	3.9125764057919437
0.01	765	3.912576405791486
0.0407	10000	2.710916520014198e+39

The first 3 end up converging to about the same value for the objective function, but 0.01 converges much faster than 0.0001. So, 0.01 was a good size in that it was small enough to coverage, but big enough to cause only a few iterations. The step size of 0.0407 was too large, and made it so the algorithm did not converge, and the value of the objective function was very large.

e.)

```
self.coef_ = np.dot(np.dot(np.linalg.pinv(np.dot(X.T, X)), X.T), y)
```

Closed Form Solution:

w = [2.44640709 -2.81635359]
cost = 3.9125764057914645
time = 0.003287792205810547 seconds

Gradient Descent:

w = [2.44640703 -2.81635347] with gradient descent, learning rate = 0.01.
cost = 3.912576405791486
time = 0.012884140014648438 seconds

Both solutions and costs are very similar; however gradient descent takes a lot longer to converge than the closed for solution.

f.)

time = 0.029319286346435547 seconds
Number of iterations: 1679
Value of objective function: 3.912576405792008
w = [2.44640678 -2.81635296]

g.)

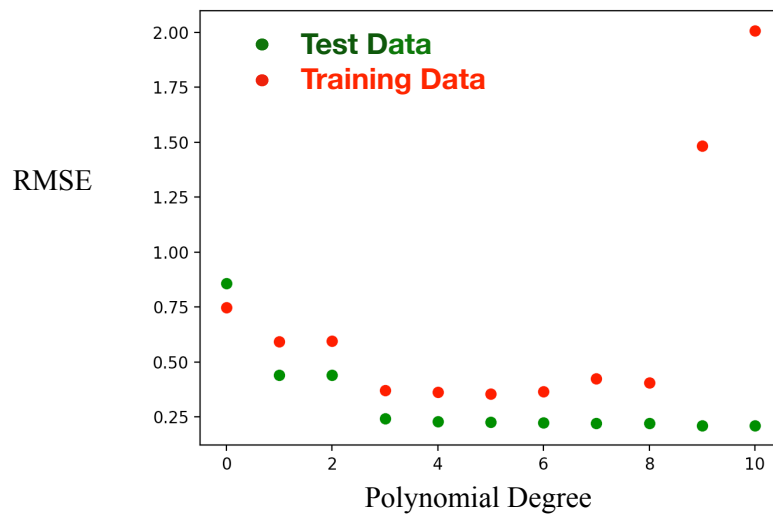
```
# part g: modify to create matrix for polynomial  
n = len(X)  
X = X.reshape(n, 1)  
Phi = np.ones((n, 1))  
m = self.m_  
for i in range (1, m+1):  
    Phi = np.append(Phi, X**i, axis=1)
```

h.)

```
# part h: compute RMSE
n, d = X.shape
error = 0
cost = self.cost(X, y)
error = (cost/n)**(1/2)
```

We prefer RMS, because it accounts for how many data points are in the set and normalizes based on that. $J(w)$ will increase as the number of points in the data set increase even if the fit is better.

i.)



I would say a polynomial of degree 5 best fits the data, because it has a low training and test root mean square error. There is evidence of under-fitting with degrees of less than 4, because both training and test error are high. There is evidence of overfitting in degrees 9 and 10, because training root mean square error is very low and testing root mean square error is very high.